

Task

Build an infinite-running agent. "Infinite-running" means the agent runs until the goal is complete - not until the context window fills. When a context window fills, the run must continue without interruption, and no information/context relevant to completing the goal is permanently lost.

What you decide

Everything - architecture, language, framework, storage, tooling.

Deliverables

1. **Working code** - runs end-to-end
2. **Demo video** - shows the agent running through a complete goal including at least two context window boundaries
3. **Approach document** - your design decisions, trade-offs, and what you noticed was missing from this spec
4. **Use-case list** - list of the use cases you tried to achieve with this agent

Evaluation

Area	What we look at
Does it work	End-to-end, including through context boundaries
System design	Are boundaries clean and extensible
Gap-filling	What did you identify as missing, and how did you handle it
Code quality	Clear boundaries, accurate naming, navigable by someone unfamiliar
Approach document	Real decisions, not post-hoc justification
Error Handling	System built must focus on production level error handling